

# COMP2006 Computer Organization

## Matrix Multiply & Time Measurement

### DGEMM

Most floating-point calculations are performed in double precision. Consider the following matrix multiply:

$$C = C + A * B$$

It is commonly called **DGEMM**, for Double precision, General Matrix Multiply. Which calculates the multiplication of matrix **A** and matrix **B** and stores the result in matrix **C**.

The following is a simple version of the **dgemm** function in the C program language:

```
void dgemm(int n, double *A, double *B, double *C)
{
    for (int i = 0; i < n; i++)
        for (int j = 0; j < n; j++)
        {
            double cij = C[i * n + j];
            for (int k = 0; k < n; k++)
                cij += A[i * n + k] * B[k * n + j];
            C[i * n + j] = cij;
        }
}
```

We use the following code to create two matrices and use the **dgemm** function to compute the results.

```
#include <stdio.h>

void dgemm(int n, double *A, double *B, double *C){...}

void main()
{
    int i, j, n = 4;
    double A[] = {1, 2, 3, 4, 2, 3, 4, 5, 3, 4, 5, 6, 4, 5, 6, 7};
    double B[] = {1, 2, 3, 4, 2, 3, 4, 5, 3, 4, 5, 6, 4, 5, 6, 7};
    double C[16] = {0};
    dgemm(n, A, B, C);
    for (i = 0; i < n; i++)
    {
        for (j = 0; j < n; j++)
            printf("%.01f\t", C[i * n + j]);
        printf("\n");
    }
}
```

You can see that we use single dimensional arrays to represent matrices.

Download **dgemm1.c** from our course web page. Then, compile the program and run it.

Compile: **gcc dgemm1.c -o dgemm1.exe**    Run: **dgemm1.exe**

You will get the results as follows:

30	40	50	60
40	54	68	82
50	68	86	104
60	82	104	126

## Time Measurement

Next, we are going to measure how many seconds the program will take for calculating the result. We use the **clock** function provided by **time.h** to measure the processing time. Therefore, we need to modify the main function as follows:

```
#include <stdio.h>
#include <sys/time.h>

void dgemm(int n, double *A, double *B, double *C){...}

int main()
{
    int i, j, n = 4;
    double A[] = {1, 2, 3, 4, 2, 3, 4, 5, 3, 4, 5, 6, 4, 5, 6, 7};
    double B[] = {1, 2, 3, 4, 2, 3, 4, 5, 3, 4, 5, 6, 4, 5, 6, 7};
    double C[16] = {0};

    struct timeval st, et;
    gettimeofday(&st, NULL);

    dgemm(n, A, B, C);

    gettimeofday(&et, NULL);
    float elapsed = ((float)(et.tv_sec - st.tv_sec))
        + (float)(et.tv_usec - st.tv_usec)*0.000001f;

    printf("Calculation time for %d x %d matrix: %0.6f seconds\n", n, n, elapsed);

    return 0;
}
```

Download **dgemm2.c** from our course web page. Then, compile the program and run it.

Compile: **gcc dgemm2.c -o dgemm2.exe**    Run: **dgemm2.exe**

You will get the results as follows:

```
Calculation time for 4 x 4 matrix: 0.000001 seconds
```

You may get different result. The calculation time is depended on the performance of your computer and the current loading of the processor.

## Dynamic Arrays

We then use some larger matrices to test the `dgemm` function. To make our program more flexible, we use dynamic array creation approach instead. We first create a double pointer and use the `malloc` function to allocate memory space with a specific size. Then, use the pointer to point to the memory space.

```
ARRAY = malloc(N * sizeof(data_type));    // create an array with N elements
```

Let's change the program as follows but we keep the `dgemm` function unchanged:

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/time.h>

void dgemm(int n, double *A, double *B, double *C){...}

void initialize(int n, double *X)
{
    int i;
    for (i = 0; i < n; i++)
        X[i] = (double)rand() / (double)RAND_MAX;
}

int main()
{
    int n;
    printf("Enter the matrix width: ");
    scanf("%d", &n);

    double *A = (double *)malloc(n * n * sizeof(double));
    double *B = (double *)malloc(n * n * sizeof(double));
    double *C = (double *)malloc(n * n * sizeof(double));

    srand(1);
    initialize(n * n, A);
    initialize(n * n, B);

    memset(C, 0, n * n * sizeof(double));

    struct timeval st, et;
    gettimeofday(&st, NULL);

    dgemm(n, A, B, C);

    gettimeofday(&et, NULL);

    float elapsed = ((float)(et.tv_sec - st.tv_sec))
        + (float)(et.tv_usec - st.tv_usec)*0.000001f;

    printf("Calculation time for %d x %d matrix: %.6f seconds\n", n, n, elapsed);

    return 0;
}
```

The `memset` statement in the program code above is used to set all elements to zero.

We also need the `initialize` function to generate some random values for the matrices. Add the following `random` function to your program file too.

Download `dgemm3.c` from our course web page. Then, compile the program and run it.

Compile: `gcc dgemm3.c -o dgemm3.exe`    Run: `dgemm3.exe`

The following is the sample output of the calculation time for two 512 x 512 matrices:

```
Enter the matrix width: 512
Calculation time for 512 x 512 matrix: 0.603715 seconds
```

*You may get different result. The calculation time is depended on the performance of your computer and the current loading of the processor.*